


Retrieval Augmented Generation (RAG)



Data saving complexities

1. Upload the documents
2. Perform the texts chunking & database saving (chunks with metadata)
3. Generate the embeddings for the saved chunk
4. Store the embeddings for the generated chunks

Ingestion

Data querying complexities

5. Ask questions, from the web interface
6. Search the question in the database, first
7. Generate embeddings for the question (if not found in the database)
8. Similarity search for the query, by embeddings → top 5 chunks, w/o ranking, w/o text search
9. LLM prompting for the user's question, with the relevant manual chunks
10. Save the question's text, embeddings and metadata
11. Return the response to the user, formatted as needed / wanted

Querying